Aditya Chandupatla

Email : aditya.chandupatla@gmail.com
https://adityachandupatla.github.io/
Mobile : 213-536-3800

EXPERIENCE

- **Coinbase** — Software Engineer
  San Francisco, California — December 2021–Current

  Data Access Control – Developed an in-house automated access control system using Apache Airflow for our Snowflake data warehouse, storing more than 10 petabytes of data. Working on extending it to be data-store-agnostic.
  Data Classification – Experimented on several classification strategies to identify and tag PII data such as: entropy measure, open source tools based on regular expressions and metadata, and third-party enterprise data-governance vendors.
  Data Retention – Designing and building Golang based microservices for Data Retention project to ensure highest compliance with regulatory requirements such as GDPR for all types of data stores across Coinbase - estimated to save the company millions of dollars.

- **Tesla** — Software Engineer and Internship
  Fremont, California — January 2021–October 2021

  Ownership Transfer Platform – Designed and implemented an orchestration engine (based on Netflix Conductor - an open-source project) for Tesla's customer account which facilitates seamless transfer of vehicles and subscriptions such as premium-connectivity, autopilot, and full-self-driving. The new asynchronous and parallel event based platform delivers 2x faster response times and supports over 2 million Tesla vehicles globally. Improved monitoring capability and visibility of the system by decoupling user-workflows resulting in 80% saved developer time.

- **Hulu** — Software Developer Intern
  Santa Monica, California — June 2020–August 2020

  Distributed Tracing System – Developed a low latency, asynchronous and secure, cloud-native ingestion service that tracks over 1 million events occurring everyday in Hulu's distributed metadata ETL pipeline. Responsible for end-to-end development, including testing, containerization, and setting up CI/CD pipeline. Deployed into production on AWS cloud using Terraform and Kubernetes. Reduced time taken by developers to search metadata documents down to single-digit seconds while not violating tight SLA requirements of pipeline.

- **Teradata** — Software Engineer
  India — July 2017–December 2018

  Analytics – Utilized micro-services architecture for automating data migration between two databases. Involved in extraction and development of feature set for time estimation task. Employed regression techniques to construct model and achieved an accuracy of 89%.
  Tensorflow – Integrated Google's distributed Tensorflow into Teradata by incorporating table-operators to provide end user with capabilities to run analytical queries right within database.
  GPU – Implemented a prototype to accelerate database aggregation operations using GPU (NVIDIA GeForce GTX 1070). Achieved a performance boost of up to 3X on a dataset containing 32 million records.

- **VMware** — Intern - IT
  India — January 2017–July 2017

  Full-stack Web Development – Took initiative to engineer utility dashboard to provide unified view of plethora of micro-services based on REST and SOAP. Led to faster deployment times, increased productivity.

EDUCATION

- **University of Southern California** — Los Angeles, CA
  Master of Science in Computer Science - 3.8/4.0 — August 2019–May 2021
  Relevant Coursework: Operating Systems, Machine Learning, Algorithms, Web Technologies

- **Jawaharlal Nehru Technological University** — India
  Bachelor of Technology in Computer Science and Engineering - 88.09% Gold Medalist — September 2013–June 2017

TECHNICAL SKILLS

- **Languages**: Java, Python, Go, C/C++, CUDA, C# (.NET), SQL, JavaScript
- **Frameworks/Tools**: Terraform, Kubernetes, Docker, Java Spring, AWS (MSK, S3, ASG's, ELB, Privatelink), Apache Kafka, Apache Airflow, Elasticsearch, Datadog, Splunk, Redis, Snowflake, Databricks, RabbitMQ